

RECENT ADVANCEMENTS IN SYMBOLIC REGRESSION AT TOP ML CONFERENCES (ICML, NEURIPS, ICLR)

HENGZHE ZHANG

VICTORIA UNIVERSITY OF WELLINGTON

06/12/2024

TABLE OF CONTENTS

1 Introduction

2 Reinforcement Learning-Based Symbolic Regression

3 Language Model-Based Symbolic Regression

4 Sparse Learning-Based Symbolic Regression

INTRODUCTION

Machine Learning for Symbolic Regression:

- Symbolic regression has garnered significant attention in recent years within the field of machine learning.
- This work focuses on leveraging neural networks to address symbolic regression problems, a topic widely discussed at top-tier ML conferences.

Categories of Deep Symbolic Regression Methods:

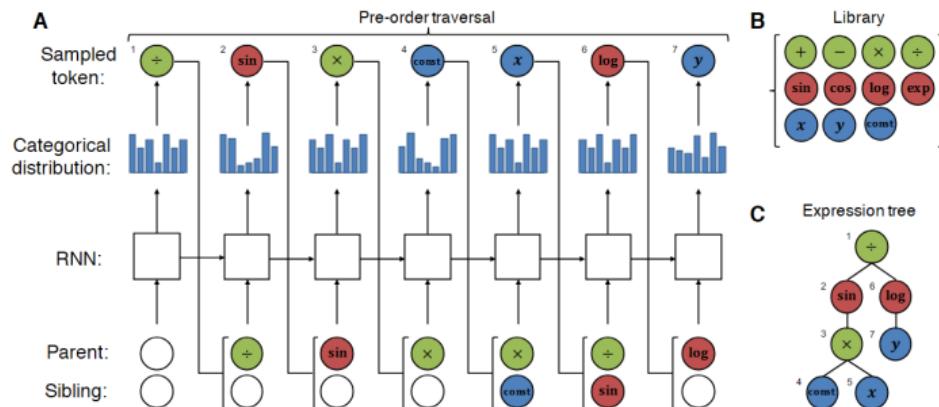
- Reinforcement Learning-Based Symbolic Regression
- Language Model-Based Symbolic Regression
- Sparse Learning-Based Symbolic Regression

REINFORCEMENT LEARNING-BASED SYMBOLIC REGRESSION

Motivation

Proposed Solution:

- Utilize a recurrent neural network (RNN) to generate mathematical expressions.
- Train using a *risk-seeking policy gradient* to optimize best-case performance.



RISK-SEEKING POLICY GRADIENT vs. STANDARD POLICY GRADIENT

Standard Policy Gradient (SPG):

- Objective:

$$J_{\text{std}}(\theta) = \mathbb{E}_{\tau \sim p(\tau|\theta)} [R(\tau)]$$

- Gradient:

$$\nabla_{\theta} J_{\text{std}}(\theta) = \mathbb{E}_{\tau \sim p(\tau|\theta)} [R(\tau) \nabla_{\theta} \log p(\tau|\theta)]$$

- Focuses on **average performance**.

Risk-Seeking Policy Gradient (RSPG):

- Objective:

$$J_{\text{risk}}(\theta; \epsilon) = \mathbb{E}_{\tau \sim p(\tau|\theta)} [R(\tau) \mid R(\tau) \geq R_{\epsilon}(\theta)]$$

- Gradient:

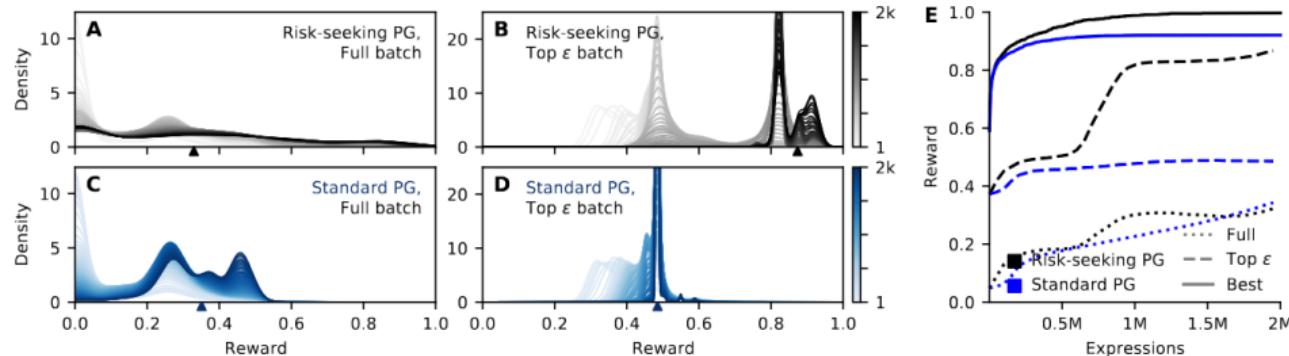
$$\nabla_{\theta} J_{\text{risk}}(\theta; \epsilon) = \mathbb{E}_{\tau \sim p(\tau|\theta)} [(R(\tau) - R_{\epsilon}(\theta)) \nabla_{\theta} \log p(\tau|\theta) \mid R(\tau) \geq R_{\epsilon}(\theta)]$$

- Focuses on **best-case performance (top- ϵ rewards)**.

Key Differences:

- RSPG modifies the gradient to emphasize high-reward trajectories, ignoring lower rewards.

RISK-SEEKING VS. STANDARD POLICY GRADIENT



■ A and B: Risk-Seeking Policy Gradient (RSPG)

- ▶ Focuses on top- ϵ rewards, prioritizing best-case outcomes.
- ▶ Shifts the distribution toward higher rewards over time.

■ C and D: Standard Policy Gradient (SPG)

- ▶ Optimizes for mean performance.
- ▶ Peaks early but lacks focus on top-reward trajectories.

■ E: Reward Comparison

- ▶ RSPG achieves higher best-case and top- ϵ rewards.
- ▶ SPG demonstrates better average reward (full batch).

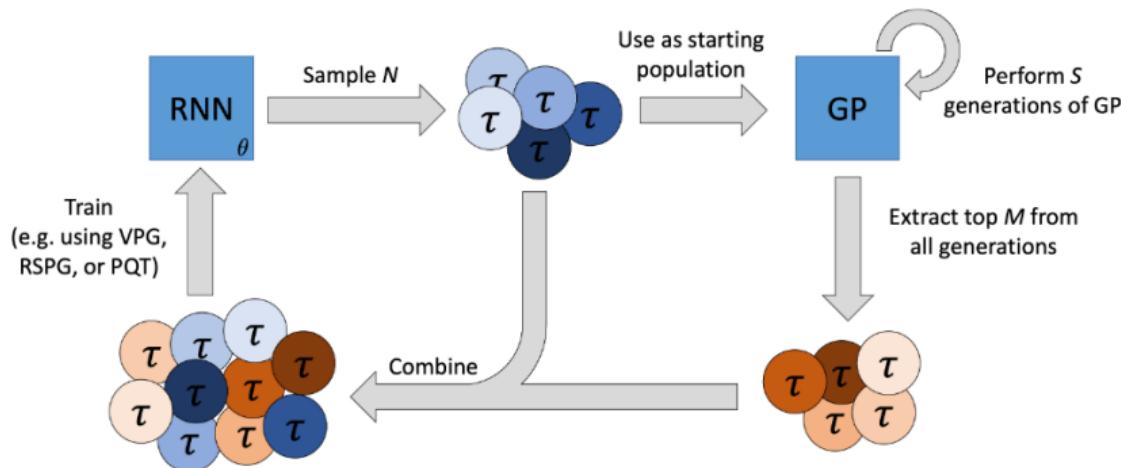
DEEP SYMBOLIC REGRESSION

Benchmark	Expression	DSR	PQT	VPG	GP	Eureqa	Wolfram
Nguyen-1	$x^3 + x^2 + x$	100%	100%	96%	100%	100%	100%
Nguyen-2	$x^4 + x^3 + x^2 + x$	100%	99%	47%	97%	100%	100%
Nguyen-3	$x^5 + x^4 + x^3 + x^2 + x$	100%	86%	4%	100%	95%	100%
Nguyen-4	$x^6 + x^5 + x^4 + x^3 + x^2 + x$	100%	93%	1%	100%	70%	100%
Nguyen-5	$\sin(x^2) \cos(x) - 1$	72%	73%	5%	45%	73%	2%
Nguyen-6	$\sin(x) + \sin(x + x^2)$	100%	98%	100%	91%	100%	1%
Nguyen-7	$\log(x + 1) + \log(x^2 + 1)$	35%	41%	3%	0%	85%	0%
Nguyen-8	\sqrt{x}	96%	21%	5%	5%	0%	71%
Nguyen-9	$\sin(x) + \sin(y^2)$	100%	100%	100%	100%	100%	-
Nguyen-10	$2 \sin(x) \cos(y)$	100%	91%	99%	76%	64%	-
Nguyen-11	x^y	100%	100%	100%	7%	100%	-
Nguyen-12	$x^4 - x^3 + \frac{1}{2}y^2 - y$	0%	0%	0%	0%	0%	-
Average		83.6%	75.2%	46.7%	60.1%	73.9%	-

SYMBOLIC REGRESSION VIA NEURAL-GUIDED GP POPULATION SEEDING (NIPS 2021)

Methodology:

- The neural network generates the initial population.
- GP evolves the population through mutation (P_m) and crossover (P_c).



SYMBOLIC REGRESSION VIA NEURAL-GUIDED GP POPULATION SEEDING (NIPS 2021)

	Recovery rate (%)			
	All	Nguyen	R	Livermore
Ours	74.92	92.33	33.33	71.09
GEGL [Ahn et al., 2020]	64.11	86.00	33.33	56.36
Random Restart GP (i.e. GP only)	63.57	88.67	2.67	58.18
DSR (i.e. RNN only) [Petersen et al., 2021]	45.19	83.58	0.00	30.41
95% confidence interval	±1.54	±1.76	±2.81	±1.32

UNIFIED DEEP SYMBOLIC REGRESSION (NIPS 2022)

Proposed Unified Framework:

- Integrates five key strategies:

- ▶ Recursive problem simplification (AI Feynman).
- ▶ Neural-guided search (Deep Symbolic Regression).
- ▶ Large-scale pre-training (NeSymReS).
- ▶ Evolutionary search (Genetic Programming).
- ▶ Sparse linear models (SINDy).

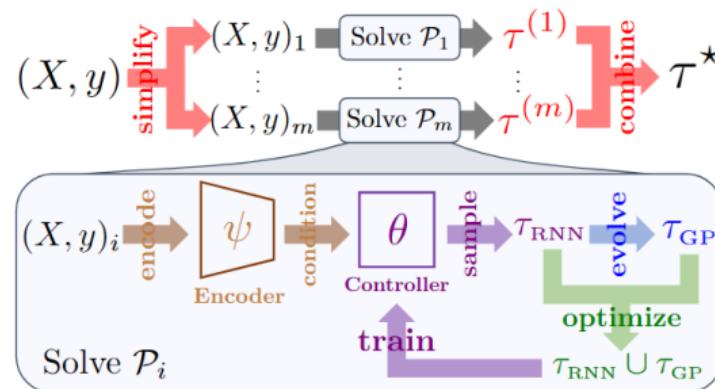
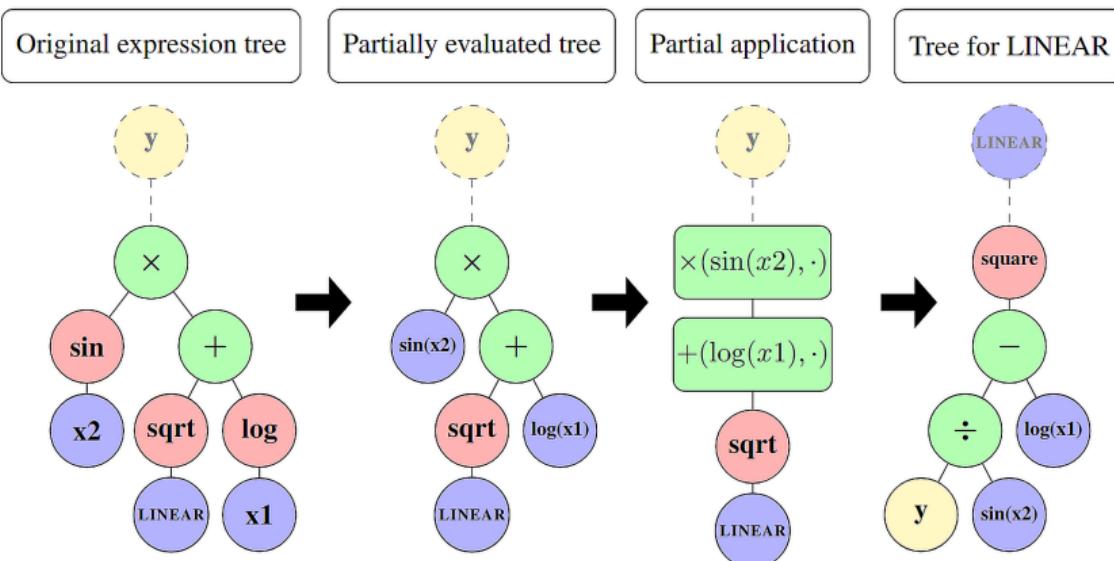
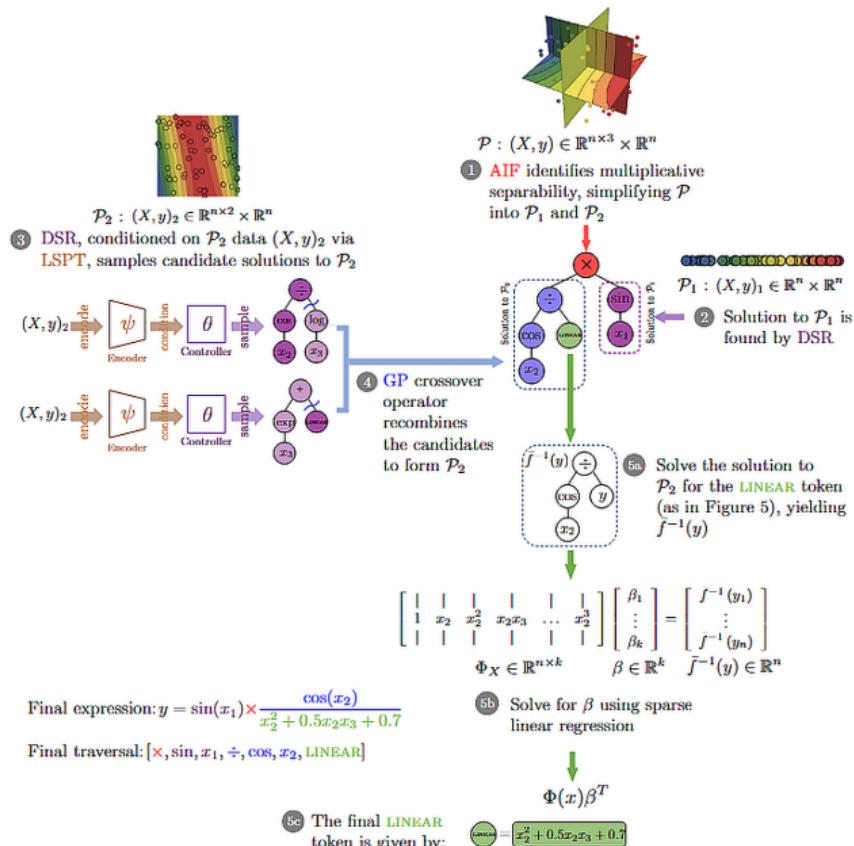


Figure 1: Unified deep symbolic regression. The five integrated solution strategies are color-coded: **AIF**, **DSR**, **LSPT**, **GP**, **LM**.

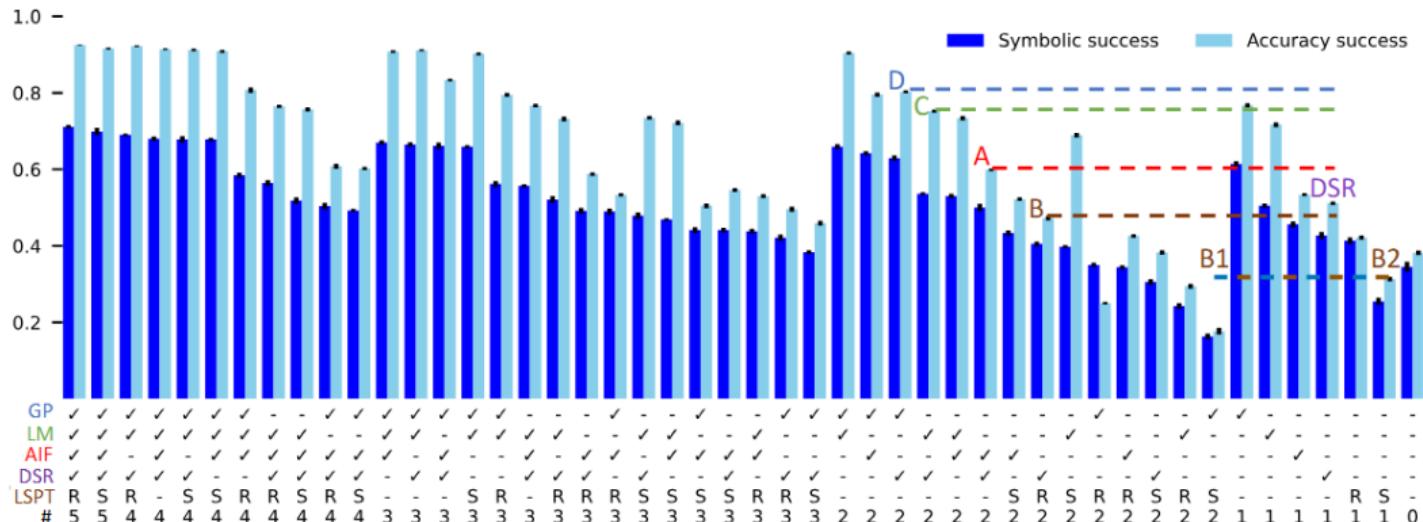
UNIFIED DEEP SYMBOLIC REGRESSION (NIPS 2022)



UNIFIED DEEP SYMBOLIC REGRESSION (NIPS 2022)



UNIFIED DEEP SYMBOLIC REGRESSION (NIPS 2022)



LANGUAGE MODEL-BASED SYMBOLIC REGRESSION

NEURAL SYMBOLIC REGRESSION THAT SCALES (ICML 2021)

Methodology:

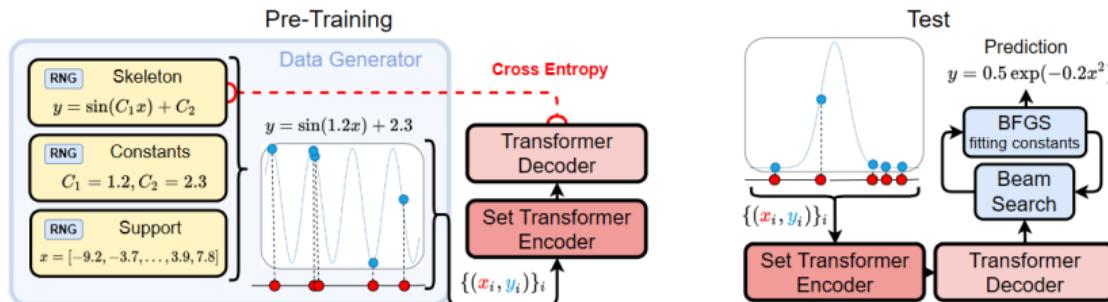
- Transformer-based architecture:

- ▶ **Encoder:** Maps $\{(x_i, y_i)\}$ to a latent space z .
- ▶ **Decoder:** Generates a skeleton e without numerical constants:

$$P(e_k|e_{<k}, z) = \text{softmax}(W_k h_k + b_k)$$

- Numerical constants are fit post-decoding using optimization methods (e.g., BFGS).
- Training uses a cross-entropy loss on skeletons:

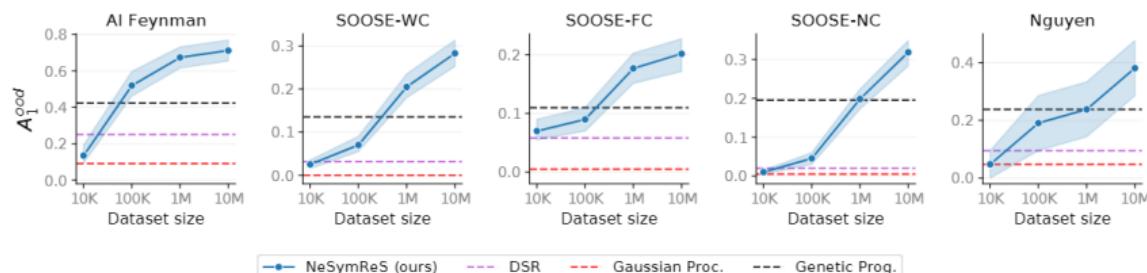
$$L = - \sum_k \log P_\theta(e_k|e_{<k}, X, Y)$$



RESULTS

Key Results:

- A_{ood} accuracy on AI-Feynman: 82.5% (NeSymReS) vs. 65.3% (Deep Symbolic Regression).
- Efficient scaling with larger datasets (up to 10M equations) and test-time data points.
- Robust against unseen equation structures (e.g., strictly out-of-sample skeletons, SOOSE).
- Limitations:
 - ▶ Supports up to $D = 3$ input dimensions.
 - ▶ Different instances of the same skeleton can represent vastly different functions depending on constants (e.g., $\sin(10x)$ vs. $\sin(0.1x)$).



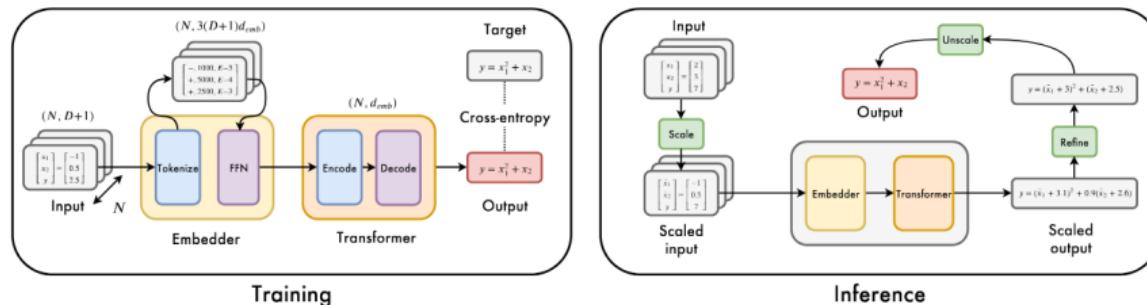
END-TO-END SYMBOLIC REGRESSION WITH TRANSFORMERS (NIPS 2022)

Methodology:

- End-to-end approach: Train Transformers to predict full expressions, including constants.
- Use symbolic-numeric tokenization:

$$f(x) = \sin(10x) \Rightarrow [\sin, *, 10, x]$$

- Refinement: Fine-tune predicted constants via BFGS with informed initialization.
- Scalable to larger input dimensions ($D = 10$) and datasets.



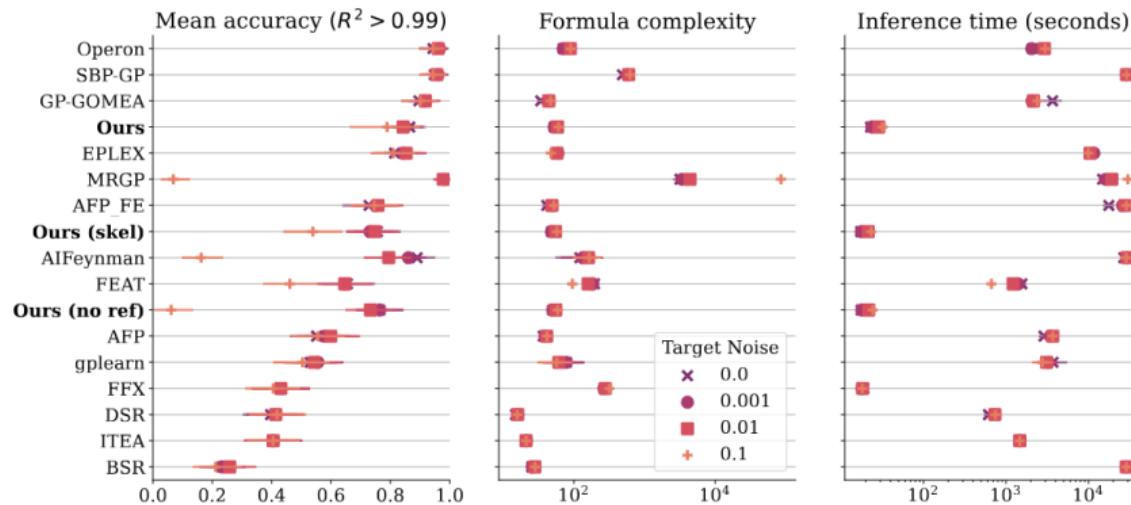
RESULTS

Key Results:

- Outperforms skeleton-based approaches in symbolic regression benchmarks:

$$R^2 = 0.68 \text{ (E2E with refinement)} \text{ vs. } R^2 = 0.43 \text{ (Skeleton-based)}$$

- Significant inference speed-up compared to genetic programming.



TRANSFORMER-BASED SYMBOLIC REGRESSION WITH JOINT SUPERVISION (ICLR 2023)

Motivation:

- Ill-posed problem: Different instances of the same skeleton behave differently.
- Poor feature extraction by previous encoders.

Proposed Method:

- **Feature Extraction:** Use a pure residual MLP framework:

$$O_i = \text{POS}(\text{MaxPool}(\text{PRE}(f_{i,j}), j \in K))$$

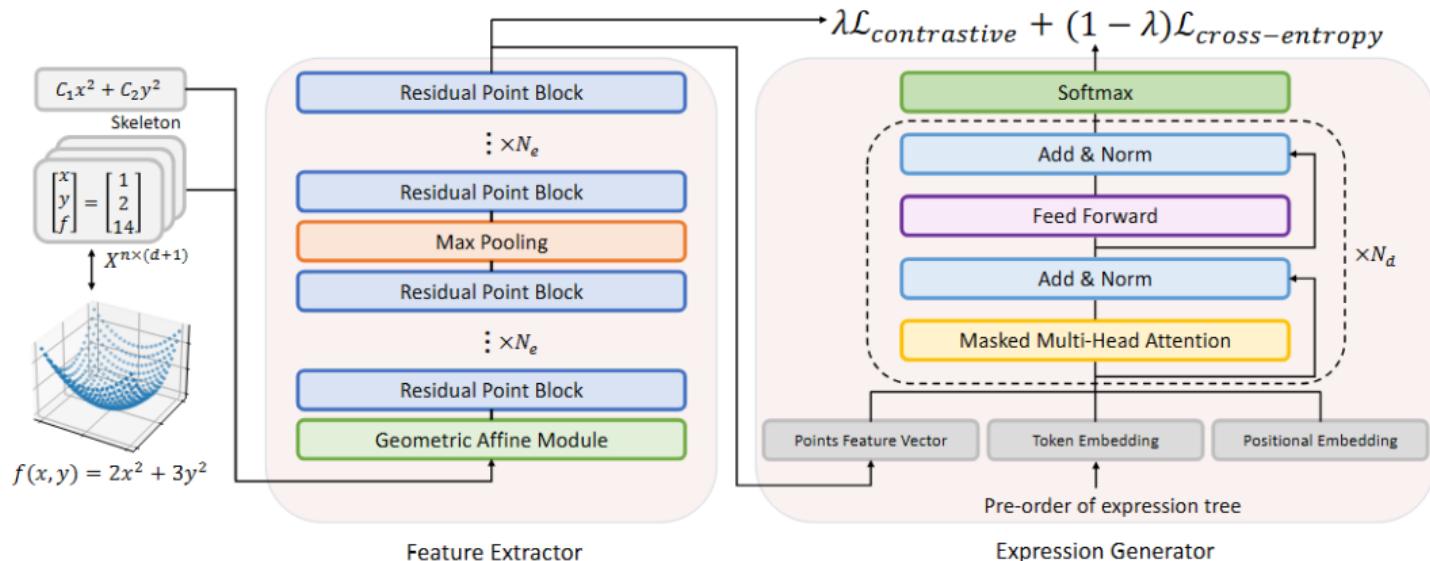
- **Joint Learning:** Combine Cross-Entropy (CE) and Contrastive Learning (CL) losses:

$$L = (1 - \lambda)L_{\text{CE}} + \lambda L_{\text{CL}}$$

- **Expression Generation:** Transformer-based autoregressive decoder:

$$P(\tau|v) = \prod_{i=1}^{|\tau|} P(\tau_i|\tau_{<i}, v)$$

TRANSFORMER-BASED SYMBOLIC REGRESSION WITH JOINT SUPERVISION



RESULTS

Benchmark	Ours	SymbolicGPT	NeSymReS	DSO	GP
	$R^2 \uparrow$				
Nguyen	0.99999	0.64394	0.97538	0.99489	0.89019
Constant	0.99998	0.69433	0.84935	0.99927	0.90842
Keijzer	0.98320	0.59457	0.97500	0.96928	0.90082
R	0.99999	0.71093	0.99993	0.97298	0.83198
AI-Feynman	0.99999	0.64682	0.99999	0.99999	0.92242
SSDNC	0.94782	0.74585	0.85792	0.93198	0.88913
Overall avg.	0.98850	0.67274	0.94292	0.97806	0.89049

LIMITATIONS OF LANGUAGE MODEL-BASED SYMBOLIC REGRESSION

Challenges:

- Symbolic regression is an **NP-hard problem**.
- Language models alone are unlikely to efficiently solve NP-hard problems.

Solutions:

- **Combine pre-trained language models with optimization techniques.**
- Transformer-Based Planning for Symbolic Regression (TPSR).

Methodology:

- Representation: Mathematical expressions are encoded as parse trees.
- MCTS Reward Function:

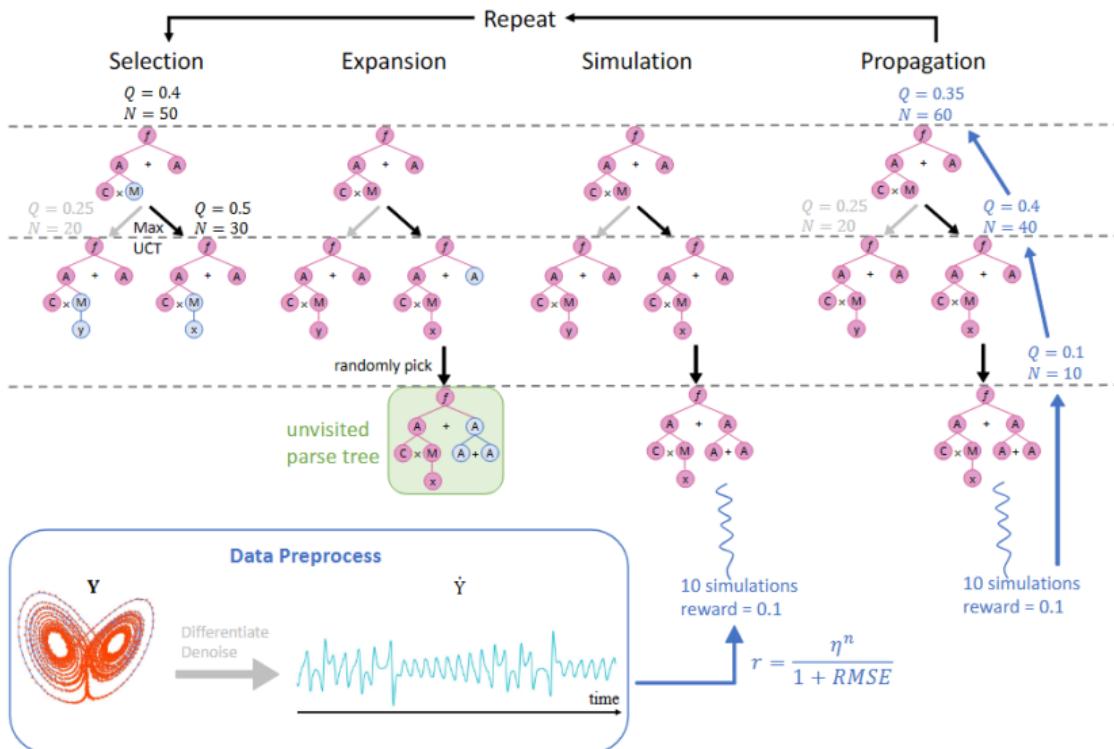
$$r = \frac{\eta^n}{1 + \sqrt{\frac{1}{N} \|Y_i - \tilde{f}(Y)\|^2}}$$

where η : parsimony factor, n : production rules count, N : number of data points.

- Upper Confidence Bound for Trees (UCT):

$$\text{UCT}(s, a) = Q(s, a) + c \sqrt{\frac{\ln N(s)}{N(s, a)}}$$

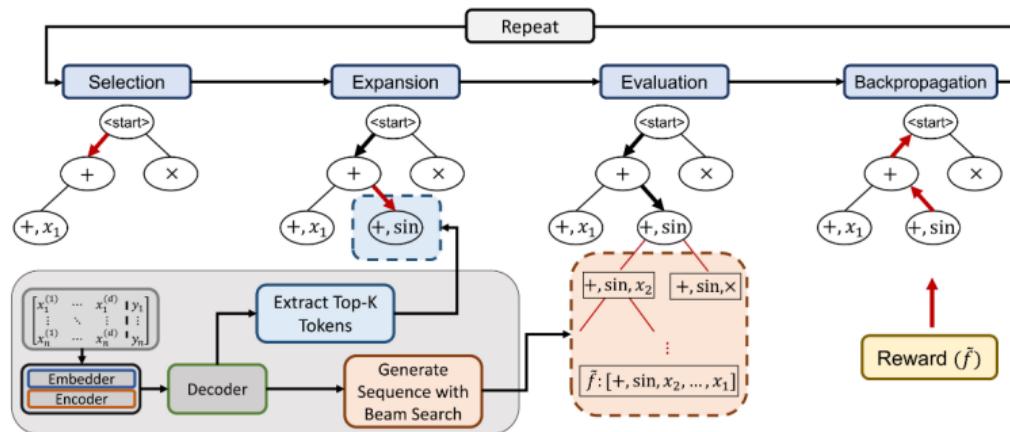
- ▶ $Q(s, a)$: Estimated reward for taking action a in state s .
- ▶ $N(s)$: Visit count for state s .
- ▶ $N(s, a)$: Visit count for action a at state s .
- ▶ c : Exploration coefficient controlling the balance between exploration and exploitation.



TRANSFORMER-BASED PLANNING FOR SYMBOLIC REGRESSION (NIPS 2023)

TPSR Objective:

- Introduce Monte Carlo Tree Search (MCTS) into Transformer decoding to guide equation generation.
- Leverage pre-trained Transformer priors while incorporating feedback on:
 - ▶ Fitting accuracy.
 - ▶ Model complexity.



Methodology:

- Transformer-based SR model enhanced with MCTS:

- ▶ s : Current state (partial equation sequence) in the Monte Carlo Tree Search (MCTS).
- ▶ a : Action (next token) extending s to a new state s' .
- ▶ $Q(s, a)$: Estimated quality (expected reward) of taking action a in state s .
- ▶ β : Hyperparameter balancing exploration and exploitation.
- ▶ $P_\theta(a|s)$: Probability of selecting action a in state s , predicted by the pre-trained Transformer model.
- ▶ $N(s)$: Visit count for state s (times s has been explored).
- ▶ $N(s')$: Visit count for the next state s' , resulting from action a .

$$P\text{-UCB}(s, a) = Q(s, a) + \beta \cdot P_\theta(a|s) \cdot \sqrt{\frac{\ln(N(s))}{1 + N(s')}}$$

$P\text{-UCB}(s, a)$

$$P\text{-UCB}(s, a) = Q(s, a) + \beta \cdot P_\theta(a|s) \cdot \sqrt{\frac{\ln(N(s))}{1 + N(s')}}$$

Key Components:

- **Exploit ($Q(s, a)$):** Prioritize high-quality actions.
- **Explore ($\sqrt{\frac{\ln(N(s))}{1 + N(s')}}$):** Encourage actions leading to less-visited states.
- **Prior ($P_\theta(a|s)$):** Leverage pre-trained Transformer knowledge to bias promising actions.

REWARD FUNCTION

Reward function combining accuracy and complexity:

$$r(f) = \frac{1}{1 + \text{NMSE}} + \lambda e^{-\frac{\ell(f)}{L}}$$

- NMSE: Normalized Mean Squared Error between predicted and actual values.
- $\ell(f)$: Length of the generated equation sequence.
- L : Maximum allowable sequence length.

SNIP: BRIDGING SYMBOLIC AND NUMERIC DOMAINS (ICLR 2024)

SNIP: A pre-training framework for joint symbolic-numeric learning using contrastive loss.

- Dual-Encoder Architecture:

$$\text{Numeric Encoder: } Z_V = A_V \cdot V_{LV}, \quad \text{Symbolic Encoder: } Z_S = A_S \cdot S_{LS}$$

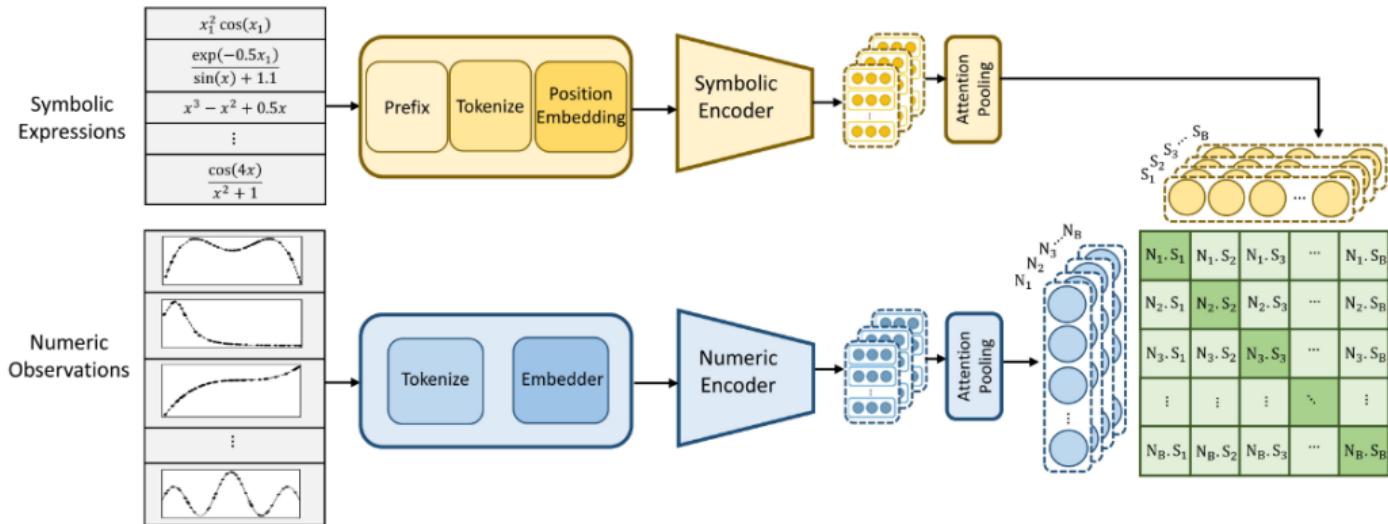
- Pre-training Objective:

$$L = - \sum_{(v,s) \in B} (\log \text{NCE}(Z_S, Z_V) + \log \text{NCE}(Z_V, Z_S))$$

where NCE aligns symbolic Z_S and numeric Z_V embeddings.

- Applications: Property prediction, symbolic regression via latent space optimization.

SNIP: BRIDGING SYMBOLIC AND NUMERIC DOMAINS (ICLR 2024)



USING SNIP FOR SYMBOLIC REGRESSION

SNIP in SR Pipeline:

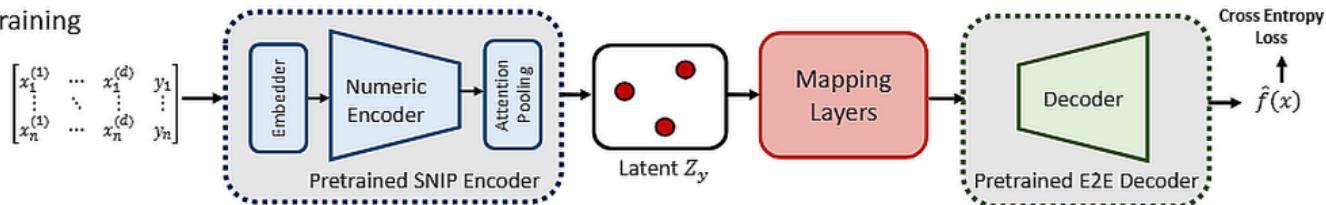
- Numeric encoder produces embeddings (Z_V).
- Integrated with a symbolic expression generation decoder (G_ω) via a mapping network.
- Training Objective: $L(f, f) = -\frac{1}{|f|} \sum_j \log P(t_j | t_1, \dots, t_{j-1}; G_\omega)$

Latent Space Optimization (LSO):

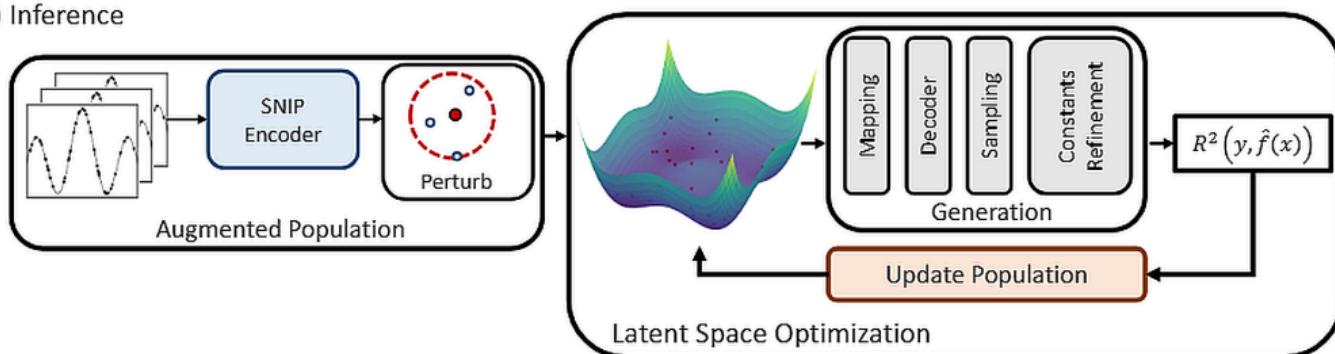
- Explore equations in the learned embedding space using optimization.
- Refine constant values for equations via numerical solvers.

SNIP: BRIDGING SYMBOLIC AND NUMERIC DOMAINS (ICLR 2024)

(a) Training



(b) Inference



SPARSE LEARNING-BASED SYMBOLIC REGRESSION

EXTRAPOLATION AND LEARNING EQUATIONS (ICLR 2017 WORKSHOP)

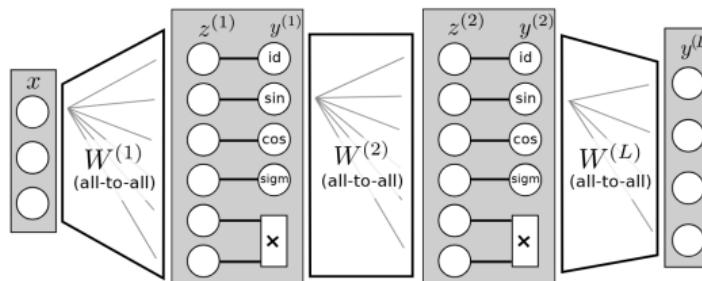
Equation Learner (EQL):

- EQL uses a feed-forward architecture with specialized units:

$$z^{(l)} = W^{(l)}y^{(l-1)} + b^{(l)}$$

$$y^{(l)} = (f_1(z_1), \dots, f_u(z_u), g_1(z_{u+1}, z_{u+2}), \dots)$$

- Unary Functions (f_i): id, sin, cos, sigmoid.
- Binary Functions (g_j): Multiplication of inputs.
- The final output is produced by a linear regression layer.



Training and Loss Function:

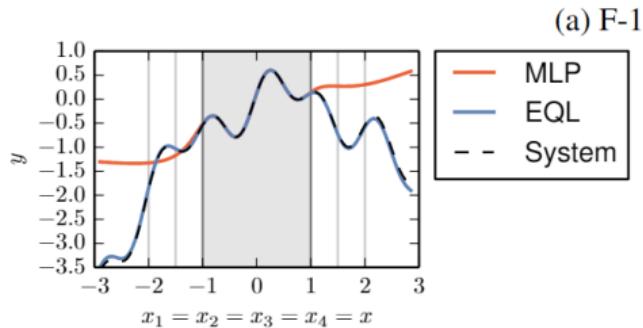
- Loss function with regularization:

$$L(D) = \frac{1}{N} \sum_{i=1}^N \|\psi(x_i) - y_i\|^2 + \lambda \sum_{l=1}^L \|W^{(l)}\|_1$$

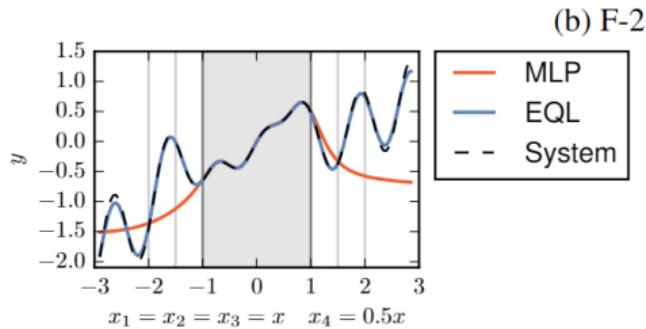
- Hybrid regularization: L_1 for sparsity and L_0 for fine-tuning.
- Regularization Strategy:
 - ▶ **Initial Phase** ($t < t_1$): No regularization ($\lambda = 0$), allowing free parameter exploration.
 - ▶ **Middle Phase** ($t_1 \leq t \leq t_2$): Apply L_1 -regularization to encourage sparsity.
 - ▶ **Final Phase** ($t > t_2$): Disable L_1 -regularization and fine-tune using L_0 -constraints.
- Transition Points:

$$t_1 = \frac{1}{4}T, \quad t_2 = \frac{19}{20}T$$

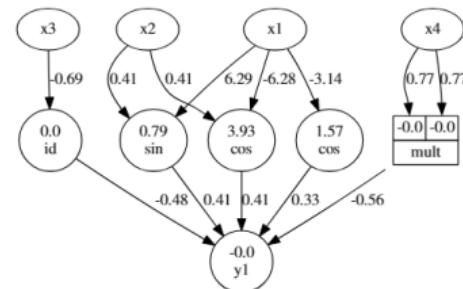
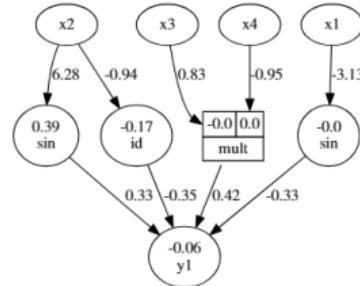
EQL RESULTS



learned formula: $-0.33 \sin(-3.13x_1) + 0.33 \sin(6.28x_2 + 0.39) + 0.33x_2 - 0.056 - 0.33x_3x_4$



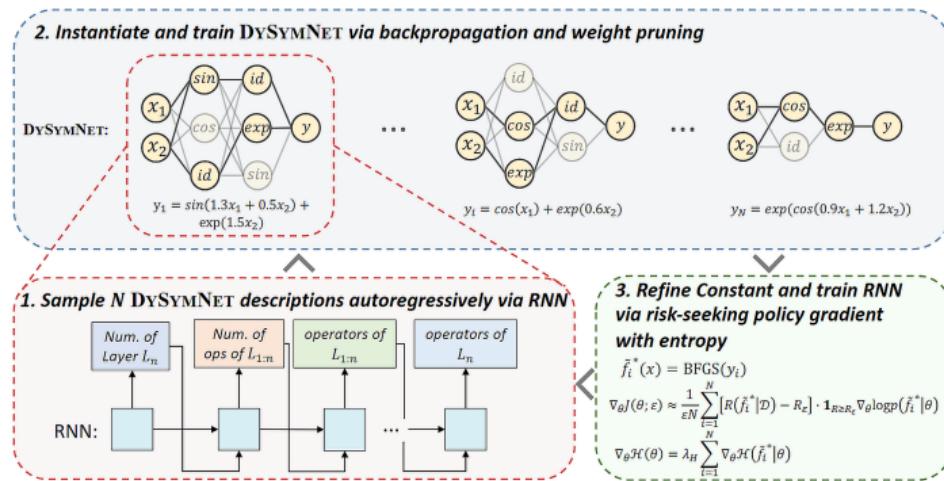
learned formula: $0.33 \cos(3.14x_1 + 1.57) + 0.33x_3 - 0.33x_4^2 + 0.41 \cos(-6.28x_1 + 3.93 + 0.41x_2) + 0.41 \sin(6.29x_1 + 0.79 + 0.41x_2)$



DYSYMNET: NEURAL-GUIDED DYNAMIC SYMBOLIC REGRESSION (ICML 2024)

- **DYSYMNET:** A novel SR framework based on reinforcement learning (policy gradient) for architecture search.
- **Policy Gradient Objective for Risk-Seeking Optimization:**

$$J_{\text{risk}}(\theta_c; \epsilon) = \mathbb{E}[R(f^*) \mid R(f^*) \geq R_\epsilon(\theta_c)].$$



DYSYMNENET RESULTS

Data Group	Benchmark	DYSYMNENET		uDSR		NGGP		NeSymReS		EQL	
		$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$
Standard $(d \leq 2)$	Nguyen	0.9999 ± 0.0001	0.9989 ± 0.0000	0.9848 ± 0.0000	0.9221 ± 0.0000	0.9758 ± 0.0032					
	Nguyen*	0.9999 ± 0.0000	0.9999 ± 0.0000	0.9999 ± 0.0001	0.3523 ± 0.0000	0.7796 ± 0.0019					
	Constant	0.9992 ± 0.0000	0.9989 ± 0.0002	0.9989 ± 0.0002	0.7670 ± 0.0000	0.9817 ± 0.0088					
	Keijzer	0.9986 ± 0.0002	0.9968 ± 0.0000	0.9912 ± 0.0005	0.9358 ± 0.0000	0.9446 ± 0.0014					
	Livermore	0.9896 ± 0.0001	0.9670 ± 0.0003	0.9665 ± 0.0005	0.9195 ± 0.0000	0.7496 ± 0.0002					
	R	0.9895 ± 0.0012	0.9999 ± 0.0000	0.9999 ± 0.0000	0.9368 ± 0.0000	0.8588 ± 0.0061					
	Jin	1.0000 ± 0.0000	0.9864 ± 0.0002	0.9942 ± 0.0001	0.9987 ± 0.0000	0.9921 ± 0.0002					
	Koza	1.0000 ± 0.0000	0.9965 ± 0.0000	0.9999 ± 0.0000	0.9999 ± 0.0000	0.7571 ± 0.0022					
Data Group	Benchmark	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$	$R^2 \uparrow$
SRBench $(d \geq 2)$	Feynman	0.9931 ± 0.0015	0.9806 ± 0.0003	0.9190 ± 0.0004	0.9234 ± 0.0000	0.5641 ± 0.0028					
	Strogatz	0.9968 ± 0.0031	0.9455 ± 0.0003	0.9534 ± 0.0003	0.8816 ± 0.0000	0.6511 ± 0.0012					
	Black-box	0.8908 ± 0.0028	0.6697 ± 0.0010	0.6086 ± 0.0021	0.4226 ± 0.0000	0.4528 ± 0.0107					

THANKS FOR LISTENING!

EMAIL: HENGZHE.ZHANG@ECS.VUW.AC.NZ

GITHUB PROJECT: [HTTPS://GITHUB.COM/HENGZHE-ZHANG/EVOLUTIONARYFOREST/](https://github.com/hengzhe-zhang/EvolutionaryForest/)